

Code Development: Best Practices

Generally and in SWIFT

```
148 *
1 * @param e The #engine.
2 * @param l A pointer to the #link, will be modified atomically.
3 * @param t The #task.
4 *
5 * @return The new #link pointer.
6 */
7 void engine_addlink(struct engine *e, struct link **l, struct task *t) {
8
9 #ifdef SWIFT_DEBUG_CHECKS
10 if (t == NULL) {
11     error("Trying to link NULL task.");
12 }
13 #endif
14
15 /* Get the next free link. */
16 const size_t ind = atomic_inc(&e->nr_links);
17 if (ind >= e->size_links) {
18     error(
19         "Link table overflow. Increase the value of "
20         "'Scheduler:links_per_tasks'.");
21 }
22 struct link *res = &e->links[ind];
23
24 /* Set it atomically. */
25 res->t = t;
26 res->next = atomic_swap(l, res);
27 }
28
29 /**
30 * @brief Repartition the cells amongst the nodes.
31 *
32 * @param e The #engine.
33 */
34 void engine_repartition(struct engine *e) {
35
36 #if defined(WITH_MPI) && (defined(HAVE_PARMETIS) || defined(HAVE_METIS))
37     ticks tic = getticks();
38
39 #endif
40 #ifdef SWIFT_DEBUG_CHECKS
```

```
188 #ifdef SWIFT_DEBUG_CHECKS
1 /* Be verbose about this. */
2 if (e->nodeID == 0 || e->verbose) message("repartitioning space");
3 fflush(stdout);
4
5 /* Check that all cells have been drifted to the current time */
6 space_check_drift_point(e->s, e->ti_current, /*check_multipoles=*/0);
7 #endif
8
9 /* Clear the repartition flag. */
10 e->forcerepart = 0;
11
12 /* Nothing to do if only using a single node. Also avoids METIS
13  * bug that doesn't handle this case well. */
14 if (e->nr_nodes == 1) return;
15
16 /* Generate the fixed costs include file. */
17 if (e->step > 3 && e->reparttype->trigger <= 1.f) {
18     task_dump_stats("partition_fixed_costs.h", e,
19         /* task_dump_threshold = */ 0.f,
20         /* header = */ 1, /* allranks = */ 1);
21 }
22
23 /* Do the repartitioning. */
24 partition_repartition(e->reparttype, e->nodeID, e->nr_nodes, e->s,
25     e->sched.tasks, e->sched.nr_tasks);
26
27 /* Partitioning requires copies of the particles, so we need to reduce the
28  * memory in use to the minimum, we can free the sorting indices and the
29  * tasks as these will be regenerated at the next rebuild. Also the foreign
30  * particle arrays can go as these will be regenerated in proxy exchange. */
31
32 /* Sorting indices. */
33 if (e->s->cells_top != NULL) space_free_cells(e->s);
34
35 /* Report the time spent in the different task categories */
36 if (e->verbose) scheduler_report_task_times(&e->sched, e->nr_threads);
37
38 /* Task arrays. */
39 scheduler_free_tasks(&e->sched);
40
```

```
505 #ifdef SWIFT_DEBUG_CHECKS
1 for (int i = 0; i < e->s->nr_cells; ++i) {
2 +--- 18 lines: const struct gravity_tensors *m = &e->s->multipoles_top[i];-----
3 }
4 #endif
5
6 /* Each node (space) has constructed its own top-level multipoles.
7  * We now need to make sure every other node has a copy of everything.
8  *
9  * We use our home-made reduction operation that simply performs a XOR
10  * operation on the multipoles. Since only local multipoles are non-zero and
11  * each multipole is only present once, the bit-by-bit XOR will
12  * create the desired result.
13  */
14 int err = MPI_Allreduce(MPI_IN_PLACE, e->s->multipoles_top, e->s->nr_cells,
15     multipole_mpi_type, multipole_mpi_reduce_op,
16     MPI_COMM_WORLD);
17 if (err != MPI_SUCCESS)
18     mpi_error(err, "Failed to all-reduce the top-level multipoles.");
19
20 #ifdef SWIFT_DEBUG_CHECKS
21 long long counter = 0;
22
23 /* Let's check that what we received makes sense */
24 for (int i = 0; i < e->s->nr_cells; ++i) {
25 +--- 13 lines: const struct gravity_tensors *m = &e->s->multipoles_top[i];-----
26 }
27 if (counter != e->total_nr_gparts)
28     error(
29         "Total particles in multipoles inconsistent with engine.\n"
30         "  counter = %lld, nr gparts = %lld",
31         counter, e->total_nr_gparts);
32 #endif
33
34 if (e->verbose)
35     message("took %.3f %s.", clocks_from_ticks(getticks() - tic),
36         clocks_getunit());
37 #else
38     error("SWIFT was not compiled with MPI support.");
39 #endif
40 }
```

Mladen Ivkovic
Durham University

Leiden, 19. September 2023

Note

- These slides contain only the part of my talk relating to the gdb/valgrind live demo.
- The state of the demos as presented then are stored in the [git repository](#) in the branch `freeze-swiftcon-2023-09`
- full slides of the workshop are available on
 - <https://1drv.ms/b/s!Aq715l3GOLnojmRg-YWc45UN9m0z?e=vabJ3l>
 - <https://mladenivkovic.github.io/work.html>

External Debugging Tools

- gdb
 - tracebacks, inspection
- valgrind
 - memory leaks
- Live demo scripts
 - <https://github.com/mladenivkovic/debugging-essentials-demo>

gdb

- Compile your program with debug symbols **-g**
- **\$ gdb** /path/to/executable
- **\$ gdb --args** path/to/executable **--arg1 --arg2**
- Useful commands:
 - **b (break)** : set a break point
 - **l (list)**: show source code in CLI
 - **p (print)**: print variables
 - For pointers: **p *pointer** works too!
 - **bt (backtrace)**: show stack trace
 - **c (continue)**: continue run
 - **pi (python-interactive)**: open a python interpreter

gdb + core dumps

- It's possible to enable core dumps:
 - when program encounters error, write down what is currently in memory instead of just quitting
 - Enabling core dumps on linux:
 - `$ ulimit -S -c unlimited`
 - Note: your sysadmins may have disabled this.
 - Default core dump location may vary. On HPC systems, it's often set to workdir. On Ubuntu 21+,
var/lib/apport/coredump/
- gdb can read that back in and allow you to debug!
- `$ gdb -c core.XXXX path/to/executable`

gdb – some tricks

- „value has been optimized out“ workaround
 - Compile entire program without optimization
 - ... or tell compiler not to optimize specific function you're looking at:

```
#pragma GCC push_options  
#pragma GCC optimize ("O0")  
your code  
#pragma GCC pop_options
```

gdb + MPI

- `$ mpirun -n 4 xterm -e gdb -ex run \`
`--args .././swift_mpi --arg1 --arg2`
 - Launch 4 MPI ranks
 - Launch a terminal (xterm) and execute subsequent command (-e)
 - Launch gdb and immediately execute „run” command
 - gdb flag: The executable (swift_mpi) will need command line arguments (--arg1, --arg2)
 - Launch swift
- *MPI_Abort()* may exit gracefully instead of raising/signalling an error.
 - Simplest workaround: Replace with *abort()* while debugging

valgrind

- Track down memory leaks
- `$ valgrind path/to/your/executable`
- `$ valgrind -leak-check=full path/to/executable`
 - For more details, traces, etc