

From `printf` to a Proper Bug Finding Strategy

Debugging and Testing Basics

Mladen Ivkovic
Durham University
November 2023



The Joy Of Finding Bugs

```
$ make
gcc main.c -o ex -g -Wall -Werror -pedantic -std=c11
19:56 [mivkov@mivkov-ThinkPad-X1-Carbon-6th] - ~/coding/d
$ ./ex
Segmentation fault (core dumped)
19:56 [mivkov@mivkov-ThinkPad-X1-Carbon-6th] - ~/coding/d
$ █
```

How to Approach Bugs

1) Don't panic.

2) Take a systematic approach. Find out:

1) What is happening?

2) Why is it happening?

3) Where is it happening?

} This is what I want
to talk about today

3) Design a unit test that reproduces the bug

4) Fix the bug

5) Document the fix (at least a comment in the code) –
your future self and collaborators will thank you!

printf – The Ol' Reliable

```
#include <stdio.h>

int main(void) {

    initialise();
    do_stuff();
    do_other_stuff();
    do_more_stuff();
    finish();

    return 0;
}
```

```
$ make
gcc main.c -o ex -g -Wall -Werror -pedantic -std=c11
19:56 [mivkov@mivkov-ThinkPad-X1-Carbon-6th] - ~/coding/d
$ ./ex
Segmentation fault (core dumped)
19:56 [mivkov@mivkov-ThinkPad-X1-Carbon-6th] - ~/coding/d
$ █
```

printf – The Ol' Reliable

Bisect the program – find out how far it gets

```
#include <stdio.h>

int main(void) {

    printf("Check 1\n");
    fflush(stdout);
    initialise();
    printf("Check 2\n");
    fflush(stdout);
    do_stuff():
    printf("Check 3\n");
    fflush(stdout);
    do_other_stuff();
    printf("Check 4\n");
    fflush(stdout);
    do_more_stuff();
    printf("Check 5\n");
    fflush(stdout);
    finish();

    return 0;
}
```

```
20:22 [mivkov@mivkov-ThinkPad-X1-Carbon-6th] - ~/codi
$ make
gcc main.c -o ex -g -Wall -Werror -pedantic -std=c11
20:22 [mivkov@mivkov-ThinkPad-X1-Carbon-6th] - ~/codi
$ ./ex
Check 1
Check 2
Check 3
Segmentation fault (core dumped)
```

WARNING: Don't forget to flush the stdout – the buffering may not display all the output!

Alternative: Set the stdbuf to unbuffered:

```
$ stdbuf -o0 ./your_program
```

printf – The Ol' Reliable

Rinse and repeat, until you find where the problem is

```
void do_other_stuff(void){  
    do_first_other_stuff();  
    do_second_other_stuff();  
    do_third_other_stuff();  
}
```



```
void do_other_stuff(void){  
    printf("In do_other_stuff: Check 1\n");  
    fflush(stdout);  
    do_first_other_stuff();  
    printf("In do_other_stuff: Check 2\n");  
    fflush(stdout);  
    do_second_other_stuff();  
    printf("In do_other_stuff: Check 3\n");  
    fflush(stdout);  
    do_third_other_stuff();  
}
```


```
20:22 [mivkov@mivkov-ThinkPad-X1-Carbon-6th] - ~/codin  
$ make  
gcc main.c -o ex -g -Wall -Werror -pedantic -std=c11  
20:28 [mivkov@mivkov-ThinkPad-X1-Carbon-6th] - ~/codin  
$ ./ex  
Check 1  
Check 2  
Check 3  
In do_other_stuff: Check 1  
In do_other_stuff: Check 2  
Segmentation fault (core dumped)  
20:28 [mivkov@mivkov-ThinkPad-X1-Carbon-6th] - ~/codin  
$
```

`printf` – The Ol' Reliable

- Bisection with printouts is good and well – but not very efficient.
 - Bisection usually means running your code over and over again until you zero in on where the program finds an error.
 - Meanwhile, the *cause* of the error may be somewhere else entirely.
 - Typically, you'd remove your checks again after you're done. So if a similar issue arises, you need to start from scratch.

Debugging Strategy

- Design your project with debugging in mind
 - Keep things clean and documented
 - Have an option to make your code talkative
 - Add tests, validations, assertions inside and outside of your code
 - Outside: Regular testing and verification.
 - Inside: On-the-fly validation, defensive programming
- Make your life easier with external tools



Let's look
at these now

Talkative Code

- Supercharge The Ol' Reliable `printf`
- Typically implemented as a runtime option determining how verbose you want your code to be.
- You can set several levels of verbosity!

Example

```
enum verbosity_level {  
    verbosity_level_quiet = 0,  
    verbosity_level_verbose = 1,  
    verbosity_level_trace = 2,  
    verbosity_level_debug = 3  
}  
  
/**  
 * Print out a message if the code is run with a high  
 * enough verbosity level.  
 * @param message Message to print out  
 * @param level threshold level of message to print.  
 * @param verbose verbosity of the run.  
 */  
void log_message(char* message,  
                | | | | | | | | verbosity_level level,  
                | | | | | | | | int verbose){  
  
    if (level <= verbose){  
        | printf("%s\n", message);  
    }  
}
```

```
int main(void) {  
  
    // Can be read in e.g. through cmdline arg  
    // or parameter file.  
    int verbose = 0;  
  
    log_message("Starting up",  
                | | | | | verbosity_level_quiet, verbose);  
    initialise();  
  
    log_message("calling do_stuff",  
                | | | | | verbosity_level_debug, verbose);  
    do_stuff();  
  
    log_message("calling do_other_stuff",  
                | | | | | verbosity_level_debug, verbose);  
    do_other_stuff();  
  
    log_message("calling do_more_stuff",  
                | | | | | verbosity_level_debug, verbose);  
    do_more_stuff();  
  
    log_message("calling finish",  
                | | | | | verbosity_level_debug, verbose);  
    finish();  
  
    log_message("Done. Bye",  
                | | | | | verbosity_level_quiet, verbose);  
  
    return 0;  
}
```

Supercharging the Logs

- You might want to consider adding:
 - Default flushing to `stdout` if verbosity level is intended for debugging
 - Writing different verbosity level outputs to different log files
 - Applying additional filters to your logs
 - All sorts of bells and whistles, e.g.
 - Which file, function, line number is this message coming from?
 - What is the elapsed time since startup?
 - What thread/rank is this message from?

External Testing

- We differentiate between
 - Unit Testing
 - Functional Testing
 - Regression Testing
- Ideally, you automate these tests.
 - e.g. every time you update your main git branch, you run a full test suite. (→ continuous integration)

Unit Testing

- Tests for small units of your code
 - e.g. individual functions, objects, ...
- Test that it does what it's supposed to, *and nothing else*
- Test corner cases, limits, how it deals with wrong data, input etc.

Unit Testing: How-To

- Split your code into logical units, and test that the units behave exactly how you want them to.
- Ideally, you have a unit test for each function you write.
- In practice, you ***need*** to have unit tests for your core routines.

Unit Testing – What to test?

- Main use case
 - Does the code do what I want it to do?
- Every corner case you *can* think of
 - Correct values, wrong values, unexpected values, ...
- Every corner case you *can't* think of
 - Try your luck here! Use randomly generated setups, values, ...
 - Advice: If you are using random numbers:
 - set a random seed to make your run reproducible
 - make your program write out the used random seed so you can reproduce the test.

Unit Testing - Example

```
/* Sorts a list of `nelems` floats stored in `array` in numerically  
 * non-decreasing order. */  
void float_sort(float array[], size_t nelems);
```

- What if `nelems == 0`?
- What if array has only one element?
- What if array has multiple identical elements?
- What if array is already sorted? Or reverse sorted?
- Does it work with an even number of elements? Odd number of elements?
- What if array contains `inf`, `-inf`, `NaN`, `FLT_MIN`, `FLT_MAX` ?

Functional Testing

- Similar to unit testing in the sense that it tests (parts of) the code
- (Parts of) code not restricted to individual units any more.
- Typically functional testing isn't concerned with the source code, but with *some external specification*.
 - E.g. *"this program must be able to read in data written in format X and create a plot."*
- In short: Test specific functionalities of (parts of) your code.

Regression Testing

- Test the actual result/output of your program for correctness.
- In a nutshell: test that known correct results are still correct.
- It's vital to have a suite of tests, experiments, benchmarks etc. to validate current state of your code.

Something Broke After I Changed It

- Strategy: Look into what you changed, and how that could have caused a problem.
- Advice: Make use of git (or any other version control)
 - If you aren't using git already, ***you really really should be***

Internal Testing

- Tests performed on-the-fly while your program runs
- Some of the trickiest bugs to uncover are due to faulty assumptions.
 - Example:
"My code has been running correctly so far."
 - Your code may end up in a faulty state without you noticing or it being detected.
 - → try to detect it!
 - Explicitly verify current state of the running program satisfies assumptions and restrictions.

Internal Testing

- Two branches of internal tests:
 - Validate the current state of your program
 - Is everything as it should be?
 - Did an unexpected situation occur?
 - Defensive programming
 - Anticipate what state could break things, and guard against them, or raise an error notifying yourself what is wrong

Validation Example

```
#include <stdio.h>

int main(void) {

    initialise();
    do_stuff();
    do_other_stuff();
    do_more_stuff();
    finish();

    return 0;
}
```

```
#include <stdio.h>

#define DEBUG_LEVEL 2

int main(void) {

    initialise();
    #if DEBUG_LEVEL > 1
        check_initialisation_corrent();
    #endif

    do_stuff();
    #if DEBUG_LEVEL > 1
        check_state_after_do_stuff();
    #endif

    // etc
    do_other_stuff();
    do_more_stuff();
    finish();

    return 0;
}
```

Validation Example

- Very often you write your validation tests while first developing a new feature.
 - ***Keep them!*** Don't throw them away.
 - If you're worried about performance, you can hide them behind a compile time macro such as

```
#if DEBUG_LEVEL > 1
```

or

```
#ifdef DEBUG_CHECKS
```
- They might automatically uncover bugs for you.
- If they don't, they will provide you with a palette of checks telling you what the problem ***isn't***. That is valuable information.

Defensive Programming

- Anticipate what state could break things, and then explicitly check that you aren't in that state.

Defensive Programming - Example

```
/**
 * @brief Returns the sound speed given density and pressure
 *
 * @param density The density \f$\rho\f$
 * @param P The pressure \f$P\f$
 * @param hydro_gamma Adiabatic index.
 */
float gas_soundspeed_from_pressure(float density,
                                   float P,
                                   float hydro_gamma) {
    return sqrtf(hydro_gamma * P / density);
}
```

$$c_s = \sqrt{\frac{\gamma P}{\rho}}$$

- What can go wrong?
 - `density == 0` → division by zero
 - Any of the variables < 0 → undefined operation
 - Any of the variables is `inf` or `NaN`

Summary So Far

- `printf` bisection does the trick, but isn't the best way of doing things.
- Design your code with debugging in mind:
 - Make it talkative, enable logging relevant information
 - Add tests and validations
 - External: Unit tests, functional tests, regression tests
 - Internal: Validation tests, defensive programming

External Debugging Tools

- Even with a good defence, bugs happen.
- Some tools to deal with them:
 - gdb
 - tracebacks, live inspection
 - Valgrind, address sanitizer
 - memory leaks, memory issues
- Live demo scripts (+ hands on exercise)
 - <https://github.com/mladenivkovic/debugging-essentials-demo>

gdb

- Compile your program with debug symbols `-g`, or `-ggdb`
 - `$ gdb /path/to/executable`
 - `$ gdb --args path/to/executable --arg1 --arg2`
- Useful commands:
 - `b (break)` set a break point
 - `l (list)` show source code in CLI
 - `p (print)` print variables
 - For pointers: `p *pointer` works too!
 - `bt (backtrace)` show stack trace
 - `c (continue)` continue run
 - `n (next)` execute next line
 - `frame` show current location in trace
 - `frame <frame nr>` change into <frame nr>

gdb breakpoints

- Set breakpoint: (*before* you execute run)
 - `break <line number in main file>`
 - e.g. `break 12` to break on line 12
 - `break path/to/file.c:<line number>`
 - e.g. `break includes/my_includes.h:13`
 - `break file.c:function_name`
 - To break when a function is called

gdb breakpoints

- Navigation:
 - `l (list)` show source code
 - `l <line number>` show source code starting at line number
 - `bt (backtrace)` show stack trace
 - `c (continue)` continue run
 - `where` show current location in trace
 - `frame` show current location in trace
 - `frame <frame nr>` change into <frame nr>
- `info break` : show information on currently set breakpoints
- `del 1` : delete breakpoint 1.
 - You can find breakpoint numbers using `info break`

gdb watchpoints

watch <variable> : halt program execution when value of variable is changed.

- You can set **watchpoints** at startup (before executing **run**) *if the variable is in global scope*.
- Otherwise, you need to first set a **breakpoint** at a point where the variable is in scope.
- Example:

```
// you can set watchpoint at  
// `sum` before `run`
```

```
int sum;  
int main(void){  
    // do stuff  
}
```

```
// you need to set a breakpoint  
// inside main() and then set a  
// watchpoint
```

```
int main(void){  
    int sum;  
    // do stuff  
}
```

gdb reverse

- It is possible to run a program in reverse with gdb, but the functionality is very restricted.
 - Many optimizations/vectorizations break it. So make sure to compile and run without any optimizations enabled.
 - Chances are it might still not work with modern instructions.
 - Recording significantly slows down execution time. May require a lot of memory.
- How-To:
 - 1) Set a **breakpoint** from which to record the program execution.
 - 2) Run the **record** command when this **breakpoint** is hit.
 - 3) Set a different **breakpoint** at the point you wish to examine.
 - 4) Once you hit the **breakpoint**, you can go back through the code using
 - **rs (reverse-step)** : go in reverse line by line
 - **rc (reverse-continue)** : go in reverse until last **breakpoint**

gdb + segfaults and other errors

- One of the most useful features of gdb is to catch signals which would usually kill your program and halt it before it aborts.
- This permits you to explore the current state of your program just before your crash.
- This doesn't need any special activation. Just run your code with gdb!
 - But don't forget to compile it with debug symbols (-g or -ggdb) to get meaningful traces

gdb + core dumps

- It's possible to enable core dumps:
 - when program encounters error, write down what is currently in memory instead of just quitting
 - Enabling core dumps on linux:
 - `$ ulimit -S -c unlimited`
 - Note: your sysadmins may have disabled this.
 - Default core dump location may vary.
 - On HPC systems, it's often set to `$workdir`.
 - On Ubuntu 21+, `/var/lib/apport/coredump/`
- gdb can read that back in and allow you to debug!
- `$ gdb -c core.XXXX path/to/executable`

gdb – “value has been optimized out”

- „value has been optimized out” workaround
 - Compile entire program without optimization
 - ... or tell compiler not to optimize specific function you’re looking at:

GCC:

```
#pragma GCC push_options
#pragma GCC optimize ("O0")

void your_function(){...}

#pragma GCC pop_options
```

Intel:

```
#pragma optimize( "", off )

void your_function() {...}

#pragma optimize( "", on )

// or

#pragma intel optimization_level 0
void your_functction(){...}
```

Clang:

```
__attribute__((optnone))
void your_function()
{...}
```

FPEs – Floating Point Exceptions

- There are ways of telling the compiler to raise an error if an invalid arithmetic operation (FPE) has occurred
- This is unfortunately not standardised – it depends on the compiler and the hardware.
- FPE types:

FE_DIVBYZERO	Pole error: division by zero, or some other asymptotically infinite result (from finite arguments).
FE_INEXACT	Inexact: the result is not exact.
FE_INVALID	Domain error: At least one of the arguments is a value for which the function is not defined.
FE_OVERFLOW	Overflow range error: The result is too large in magnitude to be represented as a value of the return type.
FE_UNDERFLOW	Underflow range error: The result is too small in magnitude to be represented as a value of the return type.

FPEs – Floating Point Exceptions

```
// `feenableexcept()` is a GNU extension, not
// standard C. We need to define _GNU_SOURCE

#define _GNU_SOURCE
#include <fenv.h>

// make sure to link with -lm as well

int main(void){

    // combine options using binary OR
    feenableexcept( FE_DIVBYZERO |
                   FE_INEXACT | FE_INVALID |
                   FE_OVERFLOW | FE_UNDERFLOW);

    // rest of your code
}
```

FPEs – Floating Point Exceptions

- With classic intel compiler:
 - Compile with
`$ icc main.c -fp-trap=mode -g`
 - (don't need to do all the `feenableexcept()` stuff)
 - Modes:

[no]divzero	Enables or disables the IEEE trap for division by zero.
[no]inexact	Enables or disables the IEEE trap for inexact result.
[no]invalid	Enables or disables the IEEE trap for invalid operation.
[no]overflow	Enables or disables the IEEE trap for overflow.
[no]underflow	Enables or disables the IEEE trap for underflow.
[no]denormal	Enables or disables the trap for denormal.
all	Enables all of the above traps.
none	Disables all of the above traps.
common	Sets the most commonly used IEEE traps: division by zero, invalid operation, and overflow.

gdb + MPI: Poor Man's Parallel Debugger

```
$ mpirun -n 4 xterm -e gdb -ex run \  
    --args your_program --arg1 --arg2
```

- Launch 4 MPI ranks
- Launch a terminal (xterm) and execute subsequent command (-e)
- Launch gdb and immediately execute „run” command
- gdb flag: The executable (your_program) will need command line arguments (--arg1, --arg2)
- Launch your program with extra cmdline arguments
- MPI_Abort() may exit gracefully instead of raising/signalling an error.
 - Simplest workaround: Replace with abort() while debugging.

valgrind

Track down memory leaks

- `$ valgrind path/to/your/executable`
- `$ valgrind --leak-check=full path/to/executable`
 - For more details, traces, etc.

Address Sanitizer

- `valgrind` is somewhat old software which might slow down your code drastically, require too much memory, or not be up-to-date with newer instructions.
- Using the address sanitizer may be a better choice.
- That is activated during compilation:
 - GCC example:
`$ gcc main.c -o ex -fsanitize=address`
 - Works the same for `clang` and `icx`.

Additional Topics

Some additional tips and tricks with gdb, which aren't covered by the live demo.

- **display <variable>** :
print variable every time it is touched throughout the run.
- **finish** :
after a **breakpoint**, run function until the end, and halt there.
- **info locals** :
print all local variables.
- **s (step)** :
like **n (next)**, but execute a single step of the line, instead of the full line.
- **disable breakpoint/watchpoint** :
disable **breakpoint** or **watchpoint**, but don't delete it. You can undo that with the **enable** command.
- **tbreak** :
set temporary **breakpoint**. Is automatically deleted after being hit once.
- **tui enable** :
start a fancy text user interface to look through program's source code. See <https://sourceware.org/gdb/current/onlinedocs/gdb.html/TUI.html> for more.

Hands-On Exercise

- Grab it from github:

<https://github.com/mladenivkovic/debugging-essentials-demo/>

The screenshot shows the GitHub repository page for 'mladenivkovic/debugging-essentials-demo'. The repository is public and has 2 branches and 3 tags. The file list shows the following files and their commit messages:

File	Commit Message	Time
demo	restructured into demo/ and exercise/	6 hours ago
exercise	removed one boundary conditions bug.	6 hours ago
slides	added slides from SWIFTcon	6 hours ago
.gitignore	added fortran version	6 hours ago
LICENSE	Initial commit	2 months ago
README	restructured into demo/ and exercise/	6 hours ago
TODO	added TODO	2 months ago
cleanup.sh	restructured into demo/ and exercise/	6 hours ago

Annotations on the image:

- A pink arrow points from the text "the demos I just went through" to the 'demo' folder.
- A green arrow points from the text "what we'll be doing now" to the 'exercise' folder.
- A pink arrow points from the text "these slides" to the 'slides' folder.

How the Game Is Played

- The `exercise/` directory contains a small toy code.
- I have planted bugs in that code.
- You go and find them!
... after a short introduction that follows

Directory Contents

- `exercise/src`
source files of toy code with planted bugs in C and in Fortran
- `exercise/solution`
correct versions of the code.
(You can compile and run this to see what the results should be.)
- `exercise/theory`
TeX documentation on what code does and how it works.
Everything you need to know about this code is on the first 2 pages!
- `exercise/plot_solution.py`
python3 script which will plot the program's output for you

The Toy Code: `linear_advection`

- Solves the equation of linear advection in 1D:

$$\frac{\partial q}{\partial t} + a \cdot \frac{\partial q}{\partial x} = 0$$


$$a = \text{const} > 0.$$

- Using a very simple numerical method:

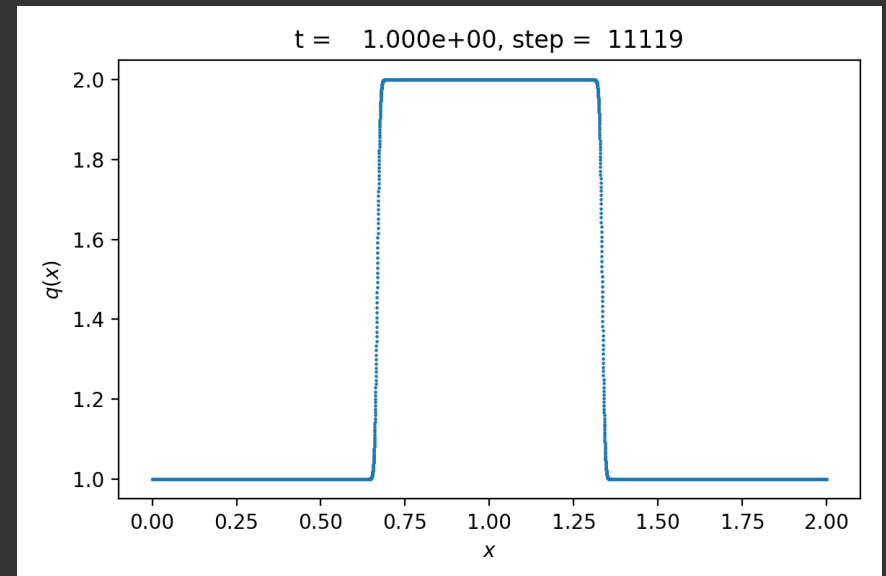



$$q_i^{n+1} = q_i^n + a \frac{\Delta t}{\Delta x} (q_{i-1}^n - q_i^n)$$

Analytical Solution

$$q(x, t) = q(x - a(t - t_0), t_0).$$




Numerical Solution



Periodic Boundary Conditions

- We add “ghost cells” along on the boundary
- Copy content of cells from boundary into ghost cells on other side



Base Algorithm of the Toy Code

- Initial Setup:
 - set $t_{\text{current}} = 0$
 - set up the initial states q_i^0 for each cell i .
- While $t_{\text{current}} < t_{\text{end}}$, solve the n-th time step:
 - Apply periodic boundary conditions
 - Find the maximal permissible time step Δt
 - For each cell i , find the updated states q_i^{n+1}
 - Update the current time: $t_{\text{current}} = t^{n+1} = t^n + \Delta t$

Determining the Time Step Size

- “CFL condition”

$$\Delta t_{max} = C_{cfl} \frac{\Delta x}{a}$$

- $C_{cfl} \in [0, 1)$: user set parameter
- Δx : cell width
- a : advection coefficient. Constant, > 0

Where to Start?

- Grab it from github:
<https://github.com/mladenivkovic/debugging-essentials-demo/>
- Check out the theory document – everything you need to know about the code is in the first 2 pages
- Run and compile the code, see what happens
- Take a look at the `Makefile` – you might want to make some changes in there
- Try using new techniques to fix it!

Instructions for DINE

- The toy program should work on your own machines. You can however also work on DINE.
 - <https://www.dur.ac.uk/icc/cosma/facilities/dine/>
 - <https://www.dur.ac.uk/icc/cosma/facilities/dine/notes/>
- Required modules:
 - A compiler:
 - Intel Classic: `module load intel_comp/2018`
 - Intel oneAPI: `module load oneAPI/2022.3.0`
 - gnu: `module load gnu_comp/10.2.0`
 - LLVM: `module load llvm/13.0.0`
 - gdb: `module load gdb`
 - valgrind: `module load valgrind`
 - WARNING: needs gnu compiler: `module load gnu_comp/10.2.0`
 - Python: `module load python/3.9.1-C8`