

# Correctness and Debugging Workshop

## Debugging Principles

*or*

### The Essentials of **gdb**

---

Mladen Ivkovic  
Durham University  
November 2025



# The Joy Of Finding Bugs

```
$ make  
gcc main.c -o ex -g -Wall -Werror -pedantic -std=c11  
19:56 [mivkov@mivkov-ThinkPad-X1-Carbon-6th] - ~/coding/de  
$ ./ex  
Segmentation fault (core dumped)  
19:56 [mivkov@mivkov-ThinkPad-X1-Carbon-6th] - ~/coding/de  
$ |
```

# printf – The Ol' Reliable

```
#include <stdio.h>

int main(void) {

    initialise();
    do_stuff();
    do_other_stuff();
    do_more_stuff();
    finish();

    return 0;
}
```

```
$ make
gcc main.c -o ex -g -Wall -Werror -pedantic -std=c11
19:56 [mivkov@mivkov-ThinkPad-X1-Carbon-6th] - ~/coding/d
$ ./ex
Segmentation fault (core dumped)
19:56 [mivkov@mivkov-ThinkPad-X1-Carbon-6th] - ~/coding/d
$ █
```

# printf – The Ol' Reliable

Bisect the program – find out how far it gets

```
#include <stdio.h>

int main(void) {
    printf("Check 1\n");
    fflush(stdout);
    initialise();
    printf("Check 2\n");
    fflush(stdout);
    do_stuff();
    printf("Check 3\n");
    fflush(stdout);
    do_other_stuff();
    printf("Check 4\n");
    fflush(stdout);
    do_more_stuff();
    printf("Check 5\n");
    fflush(stdout);
    finish();

    return 0;
}
```

```
20:22 [mivkov@mivkov-ThinkPad-X1-Carbon-6th] - ~/code
$ make
gcc main.c -o ex -g -Wall -Werror -pedantic -std=c11
20:22 [mivkov@mivkov-ThinkPad-X1-Carbon-6th] - ~/code
$ ./ex
Check 1
Check 2
Check 3
Segmentation fault (core dumped)
```

# printf – The Ol' Reliable

- Rinse and repeat, until you find where the problem is
- **WARNING:** Don't forget to **flush** the `stdout` – the buffering may not display all the output!
- Alternative: Set the `stdbuf` to unbuffered:

```
$ stdbuf -o0 ./your_program
```

# `printf` – The Ol' Reliable

- Bisection with printouts is good and well – but not very efficient.
  - Bisection usually means running your code over and over again until you zero in on where the program finds an error.
  - Meanwhile, the cause of the error may be somewhere else entirely.
  - Typically, you'd remove your checks again after you're done. So, if a similar issue arises, you need to start from scratch.

# External Debugging Tools

- There are tools which can help you to deal with them:
  - e.g. `gdb`, `DDT`, `lldb`, `mdb`, et al. for stack traces, live inspection, monitoring
  - e.g. `valgrind`, `address sanitizer` for memory leaks, memory issues
  - e.g. `MUST` for MPI issues
  - and many more

# Today's Schedule: gdb

- Live demo how to use gdb
- Hands-on exercise
- [github.com/mladenivkovic/debugging-essentials-demo](https://github.com/mladenivkovic/debugging-essentials-demo)

```
$ git clone git@github.com:mladenivkovic/debugging-essentials-demo.git
```

# github.com/mladenivkovic/debugging-essentials-demo

The screenshot shows a GitHub repository page for 'debugging-essentials-demo'. The repository is public and has 3 branches and 4 tags. The main branch is selected. There are 65 commits in total, with the most recent being 'ignore latex temp files' by Mladen Ivkovic. The commits are listed in reverse chronological order:

Commit	Message	Date
Mladen Ivkovic	ignore latex temp files	1ea560a · last year
cheatsheet	ignore latex temp files	last year
demo	cleanup notes/README's etc.	2 years ago
exercise	tested llvm support for makefiles on dine	2 years ago
slides	added slides	2 years ago

gdb cheatsheet

Examples used  
in live-demo

Hands-on  
exercises

These (and  
older) slides 9

# gdb – Basic Usage

## 1. Compilation:

Compile your program with **debug symbols**:

add `-g` or `-ggdb` flag

## 2. Running: run **gdb** executable

```
$ gdb /path/to/your/executable
```

```
$ gdb --args path/to/your/executable --arg1 --arg2
```

# gdb – Basic Commands

r (run)	run the program
b (break)	set a break point
l (list)	show source code in CLI
p (print)	print variables. For pointers: p *pointer works too!
bt (backtrace)	show stack trace
c (continue)	continue run (e.g. after hitting breakpoint, halting...)
n (next)	execute next line
frame or where	show current location in trace
frame <frame nr>	change into <frame nr>

# gdb – Breakpoints

Set breakpoint: (*before* you execute run)

`break <line number in main file>`

e.g. `break 12` to break on line 12

`break path/to/file.c:<line number>`

e.g. `break includes/my_includes.h:13`

`break file.c:function_name`

To break when a function is called

# gdb – Breakpoints Navigation

l (list)	show source code
l <line number>	show source code at line number
bt (backtrace)	show stack trace
c (continue)	continue run
where	show current location in trace
frame	show current location in trace
frame <frame nr>	change into <frame nr>

# gdb – More on Breakpoints

`info break` : show information on currently set breakpoints

`del 1` : delete breakpoint 1.

You can find breakpoint numbers using `info break`

`break <location> if <condition>`

Set a **conditional breakpoint** at location (e.g. line number, function) and break only if condition is satisfied, e.g.

`break main.c:12 if i == 12 && j < 23`

# gdb – Watchpoints

`watch <variable>` : halt program execution when value of variable is changed.

You can set `watchpoints` at startup (before executing `run`) *if the variable is in global scope.*

Otherwise, you need to first set a `breakpoint` at a point where the variable is in scope.

# gdb – Watchpoints Scope Example

Variable `sum` in global scope

```
// you can set watchpoint  
// at `sum` before `run`  
  
int sum;  
  
int main(void){  
    // do stuff  
}
```

Variable `sum` **not** in global scope

```
// you need to set a  
// breakpoint inside main()  
// and then set a  
// watchpoint for `sum`  
  
int main(void){  
    int sum;  
    // do stuff  
}
```

## gdb - reverse

- It is possible to run a program **in reverse** with **gdb**, but the functionality is very restricted.
- Many optimizations/vectorizations break it. So make sure to *compile and run without any optimizations* enabled.
- Chances are it might still not work with modern instructions.
- Recording *significantly slows down* execution time. May require a lot of memory.

# gdb – reverse: How-To

- 1) Set a **breakpoint** from which to record the program execution.
- 2) Run the **record** command when this breakpoint is hit.
- 3) Set a different **breakpoint** at the point you wish to examine.
- 4) Once you hit the second **breakpoint**, you can go back through the code using

**rs** (reverse-step)

go in reverse line by line

**rc** (reverse-continue)

go in reverse until last breakpoint

# gdb – Segfaults and Other Errors

- One of the most useful features of `gdb` is to **catch signals** which would usually kill your program and halt it before it aborts.
- This permits you to **explore the current state** of your program just before your crash.
- This doesn't need any special activation. Just run your code with `gdb`!
  - *But don't forget to compile it with debug symbols (`-g` or `-ggdb`) to get meaningful traces!*

# gdb – Core Dumps

When running code, it is possible to enable core dumps:

when program encounters error, write down what is currently in memory instead of just quitting

Typically writes `core.XXXX` files, where `XXXX` are some numbers.

gdb can read that back in and allow you to debug!

```
$ gdb -c core.XXXX path/to/your/executable
```

# gdb – Core Dumps: How to Enable and Where to Find

Enabling core dumps on linux:

```
$ ulimit -S -c unlimited
```

**Note:** your sysadmins may have disabled core dumps.

Default core dump location may vary:

On HPC systems, it's often set to \$workdir

Ubuntu 21+: /var/lib/apport/coredump/

Manjaro: /var/lib/systemd/coredump/

# gdb – Core Dumps With `coredumpctl`

```
$ coredumpctl info
```

Show info of last core dump.

To find out where it's stored, look for line starting with "Storage"

```
$ coredumpctl debug
```

Launch the debugger on the last available core dump

```
$ coredumpctl debug --debugger=/path/to/gdb
```

Specify which debugger to use

```
$ coredumpctl dump --output core.out
```

```
$ gdb -c core.out /path/to/your/executable
```

Some operating systems will compress their core dumps, making them unreadable by `gdb`. These two lines uncompress the last core dump into `core.out` and read them in using `gdb`

# gdb – “value has been optimized out”

„value has been optimized out“ workarounds:

- Compile entire program without optimisation: `-O0` compiler flag  
(You may be able to recompile only a handful of files you’re interested in using that flag, leaving the rest optimised)
- Or tell compiler not to optimize specific function you’re looking at:

**GCC:**

```
#pragma GCC push_options
#pragma GCC optimize ("O0")

void your_function(){...}

#pragma GCC pop_options
```

**Intel:**

```
#pragma optimize( "", off )
void your_function() {...}
#pragma optimize( "", on )

// ---- or, alternatively ----

#pragma intel optimization_level 0
void your_funtction() {...}
```

**Clang:**

```
--attribute__(optnone)
void your_function()
{...}
```

# FPEs – Floating Point Exceptions

- There are ways of telling the compiler to **raise an error if an invalid arithmetic operation (FPE)** has occurred
- This is unfortunately not standardised – it depends on the compiler and the hardware.

# FPEs – Types of Floating Point Exceptions

FE_DIVBYZERO	Pole error: division by zero, or some other asymptotically infinite result (from finite arguments).
FE_INEXACT	Inexact: the result is not exact.
FE_INVALID	Domain error: At least one of the arguments is a value for which the function is not defined.
FE_OVERFLOW	Overflow range error: The result is too large in magnitude to be represented as a value of the return type.
FE_UNDERFLOW	Underflow range error: The result is too small in magnitude to be represented as a value of the return type.

# FPEs – Floating Point Exceptions in C

```
// `feenableexcept()` is a GNU extension, not
// standard C. We need to define _GNU_SOURCE

#define _GNU_SOURCE
#include <fenv.h>

// make sure to link with -lm as well

int main(void){

    // combine options using binary OR
    feenableexcept(FE_DIVBYZERO | FE_INEXACT | FE_INVALID |
                   FE_OVERFLOW | FE_UNDERFLOW);

    // rest of your code
}
```

# gdb + MPI: Poor Man's Parallel Debugger

- Running `gdb` with `MPI`: Simultaneously launch one `gdb` instance per `MPI` rank in a terminal emulator (`xterm`)

```
$ mpirun -n 4 xterm -e gdb -ex run --args your_program --arg1 --arg2
```

- *Launch 4 MPI ranks*
- *Launch a terminal (xterm) and execute subsequent command (-e)*
- *Launch gdb and immediately execute „run” command*
- *gdb flag: The executable (your\_program) will need command line arguments  
(--arg1, --arg2)*
- *Launch your program with extra cmdline arguments*

# `gdb` + MPI: Poor Man's Parallel Debugger

- **Warning:** depending on the implementaiton, `MPI_Abort()` may exit gracefully instead of raising/signalling an error, killing the `gdb` session instead of letting it catch and handle the abort.
- Simplest workaround: Replace `MPI_Abort()` with `abort()` while debugging.

# gdb – Additional Topics, Tricks etc

- `display <variable>`  
print variable every time it is touched throughout the run.
- `finish`  
after a breakpoint, run function until the end, and halt there.
- `info locals`  
print all local variables.
- `s (step)`  
like `n (next)`, but execute a single step of the line, instead of the full line.
- `disable <breakpoint-nr>/<watchpoint-nr>`:  
disable `breakpoint` or `watchpoint` by number (use `info break` or `info watch` to find number), but don't delete it.  
You can undo that with the `enable` command.
- `tbreak`  
set temporary breakpoint: will be automatically deleted after being hit once.

# gdb – Additional Topics, Tricks etc

- `tui enable`:  
start a fancy text user interface to look through program's source code.  
See <https://sourceware.org/gdb/current/onlinedocs/gdb.html/TUI.html> for more.
- `make`: run `make` from within `gdb`
- `shell`: open a shell (`bash/zsh/csh` etc, your default) inside `gdb`
- `python`: open a `python` interpreter within `gdb`
- `save breakpoints <filename>`:  
Save `breakpoints` of current session into `<filename>`
- `source <filename>`:  
Load (e.g. saved `breakpoints`) from `<filename>`

# Hands-On Exercise

```
$ git clone git@github.com:mladenivkovic/debugging-essentials-demo.git
```

commit	Author	Message	Date	Commits
1ea560a · last year	Mladen Ivkovic	ignore latex temp files	last year	65 Commits
cheatsheet		ignore latex temp files		last year
demo		cleanup notes/README's etc.		2 years ago
exercise		tested llvm support for makefiles on dine		2 years ago
slides		added slides		2 years ago

gdb cheatsheet

Examples used  
in live-demo

Hands-on  
exercises

These (and  
older) slides 31

# How the Game Is Played

- The `exercise/` directory contains a small toy code.
- I have planted bugs in that code.
- You go and find them!

... after a short introduction that follows

# Directory Contents

- **exercise/src**
  - source files of toy code with planted bugs in C and in Fortran
- **exercise/solution**
  - correct versions of the code.
  - (You can compile and run this to see what the results should be.)
- **exercise/theory**
  - TeX documentation on what code does and how it works.
  - Everything you need to know about this code is on the first 2 pages!
- **exercise/plot\_solution.py**
  - `python3` script which will plot the program's output for you

# The Toy Code: linear\_advection

- Solves the equation of linear advection in 1D:

$$\frac{\partial q}{\partial t} + a \cdot \frac{\partial q}{\partial x} = 0 \quad a = \text{const} > 0.$$


- Using a very simple numerical method:

$$q_i^{n+1} = q_i^n + a \frac{\Delta t}{\Delta x} (q_{i-1}^n - q_i^n)$$





# Analytical Solution

- Original solution at  $t = t_0$  moved (adverted) by  $a(t - t_0)$

$$q(x, t) = q(x - a(t - t_0), t_0).$$




# Numerical Solution



# Periodic Boundary Conditions

- We add “ghost cells” along on the boundary
- Copy content of cells from boundary into **ghost cells** on other side



# Base Algorithm of the Toy Code

Initial Setup:

set  $t_{current} = 0$

set up the initial states  $q_i^0$  for each cell  $i$ .

While  $t_{current} < t_{end}$ , solve the n-th time step:

Apply periodic boundary conditions

Find the maximal permissible time step  $\Delta t$

For each cell  $i$ , find the updated states  $q_i^{n+1}$

Update the current time:  $t_{current} = t^{n+1} = t^n + \Delta t$

# Determining the Time Step Size

- “CFL condition”

$$\Delta t_{max} = C_{cfl} \frac{\Delta x}{a}$$

- $C_{cfl} \in [0, 1)$  : user set parameter
- $\Delta x$  : cell width
- $a$ : advection coefficient. Constant, positive, non-zero.

# Where to Start?

- Grab it from github:  
<https://github.com/mladenivkovic/debugging-essentials-demo/>
- \$ git clone git@github.com:mladenivkovic/debugging-essentials-demo.git
- Check out the theory document – everything you need to know about the code is in the first 2 pages
  - Compile and run the code, see what happens
  - Take a look at the **Makefile** – you might want to make some changes in there
  - Try using **gdb** to fix it!

# Final Slide

# Backup Slides

# valgrind

Track down memory leaks

```
$ valgrind path/to/your/executable
```

```
$ valgrind --leak-check=full path/to/executable
```

For more details, traces, etc.

# Address Sanitizer

- `valgrind` is somewhat old software which might slow down your code drastically, require too much memory, or not be up-to-date with newer instructions.
- Using the `address sanitizer` may be a better choice.
- That is activated during compilation:

GCC example:

```
$ gcc main.c -o ex -fsanitize=address
```

Works the same for `clang` and `icx`.

- ***Note:*** for more complex compilations `-fsanitize=address` must also be passed to the linker!